



# Greedy Algorithms

- The simplest class of algorithms are **greedy algorithms**.
  - They make a *locally optimal* (or *myopic*) choice at each step.
- Greedy algorithms have several nice properties:
  - *Fast.*
  - *Simple.*
  - *Easy to Code.*
- Unfortunately, they are also usually rubbish.



So for this topic the goal is to understand:

- *The basic greedy algorithmic techniques.*
- *On which problems these techniques work.*

# Task Scheduling



- A firm receives job orders from  $n$  customers.
  - The firm can do only one task at a time.
  - The job of customer  $i$  will take the firm  $t_i$  units of time to complete.
  - Each customer  $i$  wants their job completed as *early as possible* (and will pay accordingly).
- Assume the firm processes the jobs in the order  $\{1, 2, 3, \dots, n\}$ .
  - Then the **waiting time** of customer  $i$  will be:  $w_\ell = \sum_{i=1}^{\ell} t_i$
- So to maximize profits the objective of the firm is to *minimize the sum of the waiting times*:

$$\sum_{\ell=1}^n w_\ell = \sum_{\ell=1}^n \sum_{i=1}^{\ell} t_i$$



## Greedy Task Scheduling

- Here is a **greedy algorithm** that solves this problem optimally.

### Greedy Task Scheduling Algorithm

Sort the jobs by length  $t_1 \leq t_2 \leq \dots \leq t_n$

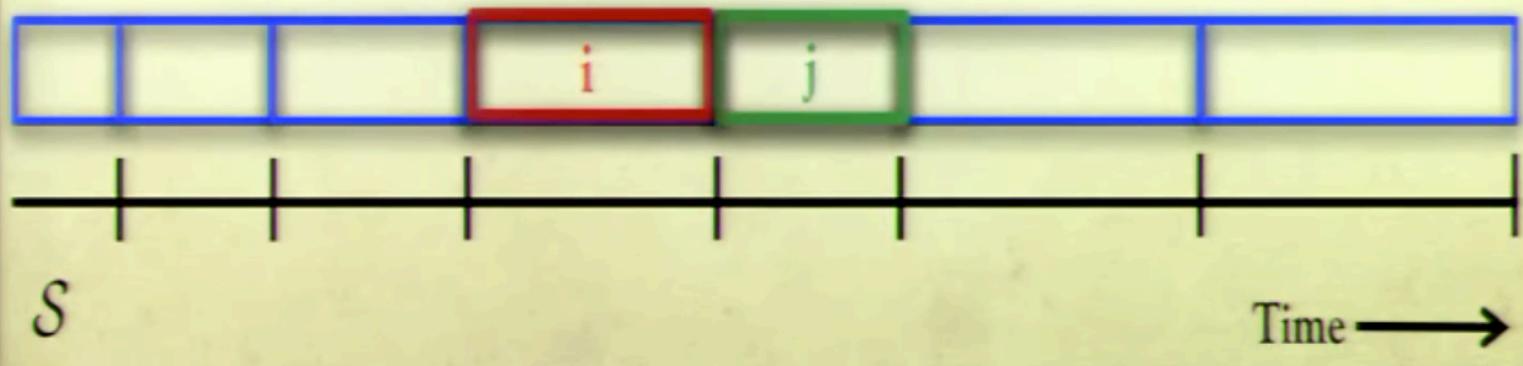
Schedule the jobs in the order  $\{1, 2, 3, \dots, n\}$

# An Exchange Argument

**Theorem.** The greedy algorithm outputs an *optimal* schedule.

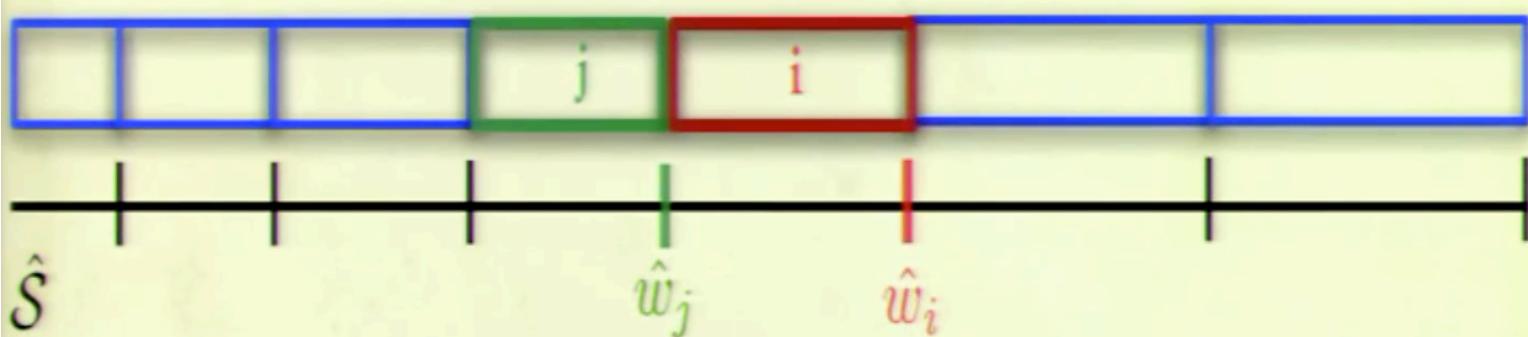
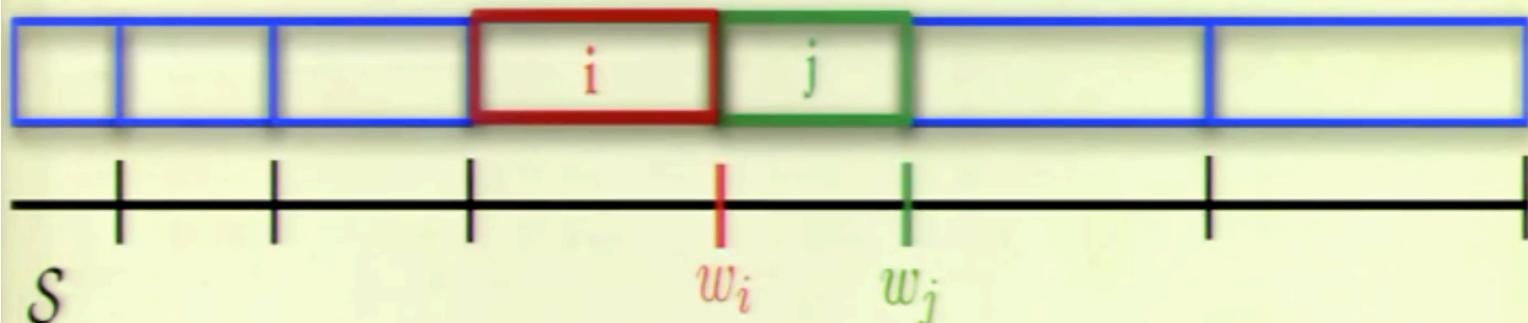
**Proof.**

- Let the greedy algorithm schedule in the order  $\{1, 2, \dots, n\}$
- Assume there is a better schedule  $S$
- Then there must be a pair of jobs  $i$  and  $j$  such that:
  - Job  $i$  is scheduled immediately before job  $j$  by schedule  $S$
  - Job  $i$  is longer than job  $j$ :  $t_i > t_j$



## Proof [cont.]

- But then exchanging the order of jobs  $i$  and  $j$  gives a better schedule  $\hat{S}$ .



- Observe the *waiting time* of every other job remains the same:

$$\hat{w}_k = w_k \quad \forall k \neq i, j$$

- But clearly:  $\hat{w}_i + \hat{w}_j < w_i + w_j$
- This contradicts the assumption that  $S$  was an optimal schedule. □

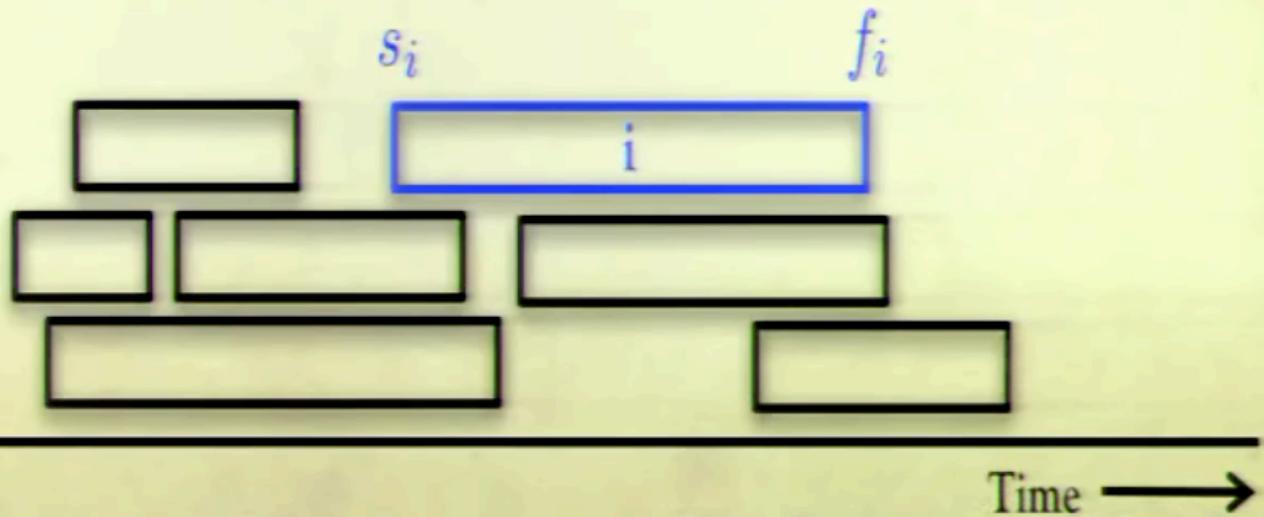
# The Running Time



- How long does the greedy task scheduling algorithm take?
  - We need only sort  $n$  time lengths.
$$\implies \text{Runtime} = O(n \cdot \log n)$$
- So this algorithm is efficient.

# Class Scheduling

- There is one classroom.
- There is a set  $I = \{1, 2, \dots, n\}$  of classes that request the room.
  - Class  $i$  has a start time  $s_i$  and a finish time  $f_i$ .
- The objective is to book as many classes into the room as possible.
  - However, we **cannot** book two classes that need to use the room at exactly the *same time*.
- This problem is often called the **Interval Selection Problem**, as we can plot the classes as intervals between their start and finish times.





# Greedy Scheduling Algorithms

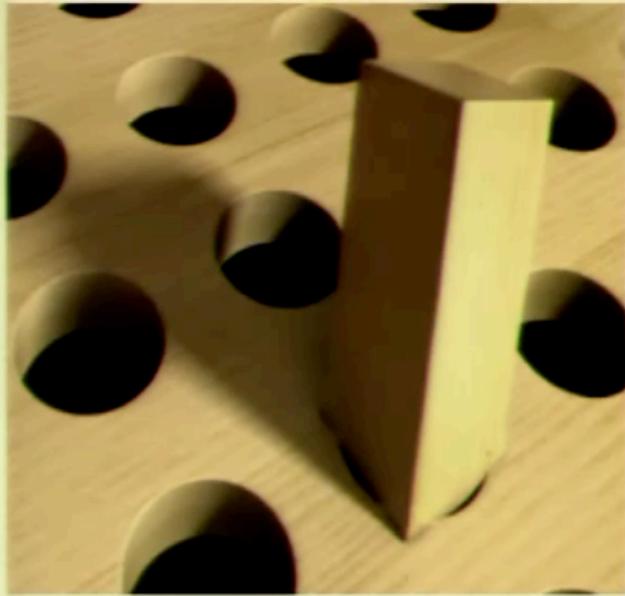
- But which greedy algorithm? *There are many possibilities:*

**First-Start.** Select the class that starts earliest, and iterate on the remaining classes that do not conflict with this selection.

**Shortest-Duration.** Select the class of shortest duration, and iterate on the remaining classes that do not conflict with this selection.

**Minimum-Conflict.** Select the class that conflicts with the fewest number of classes, and then iterate...

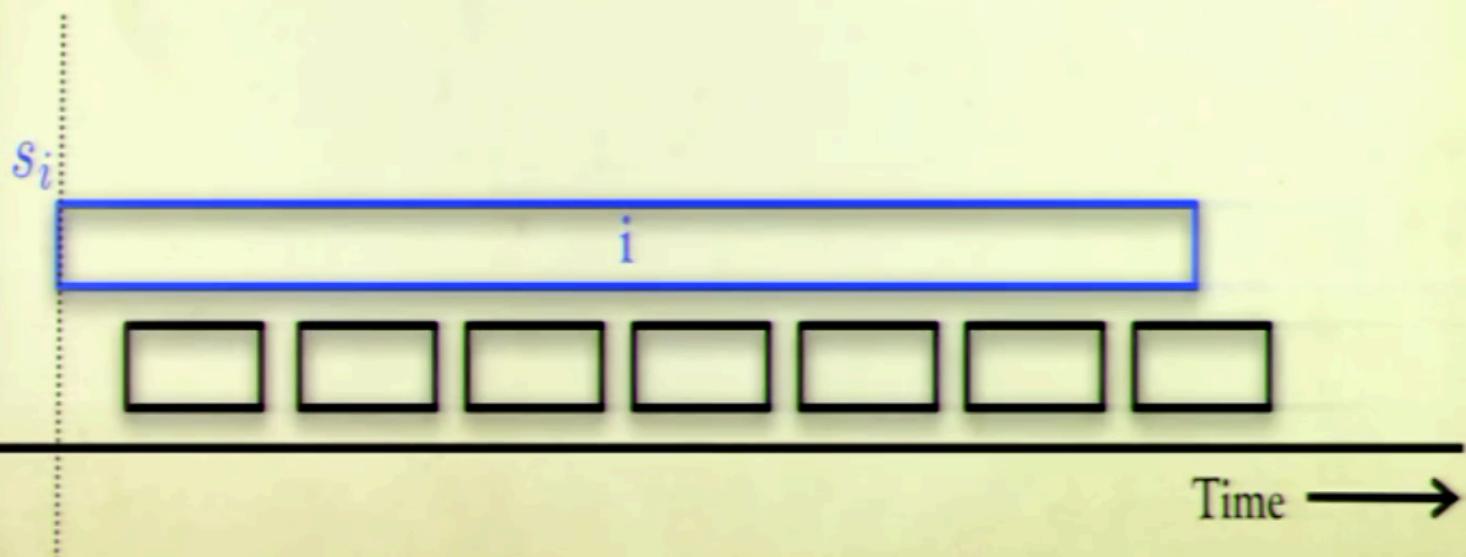
**Last-Start.** Select the class that starts last, and then iterate...

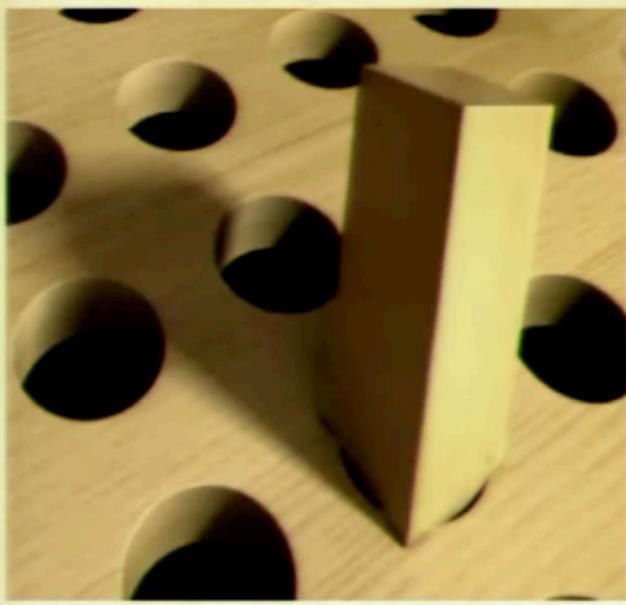


## First-Start

- Do these any of these algorithms even work?

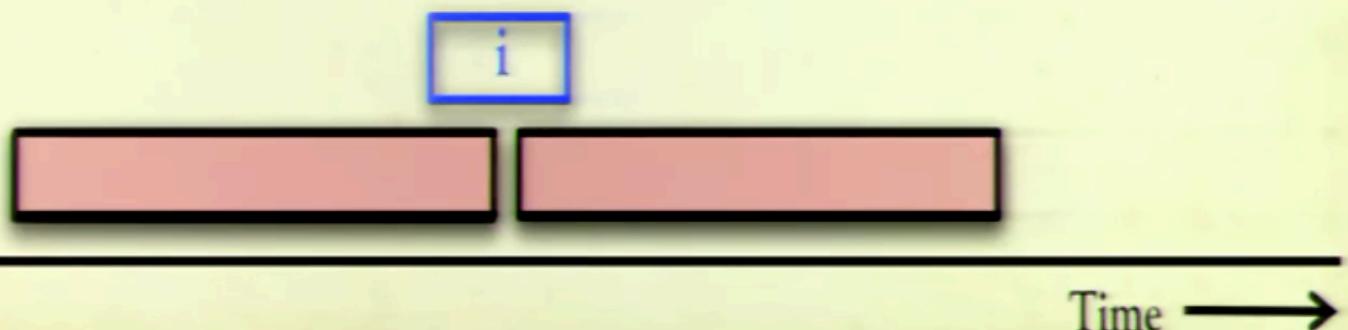
**First-Start.** Select the class that starts earliest, and iterate on the remaining classes that do not conflict with this selection.

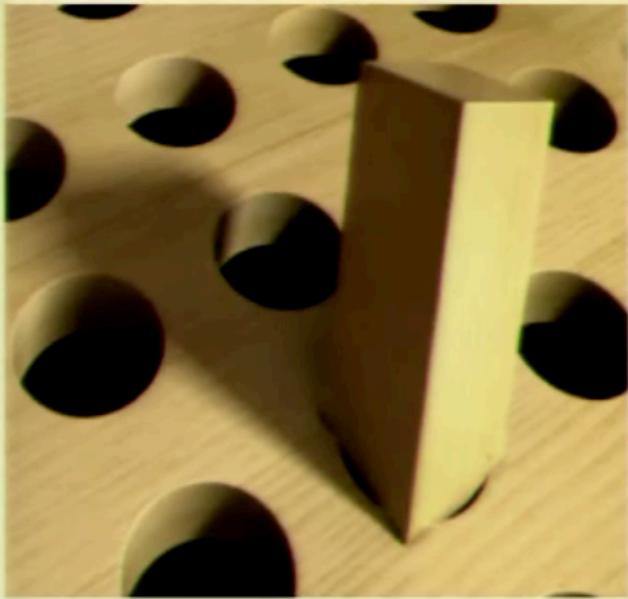




## Shortest-Duration

**Shortest-Duration.** Select the class of shortest duration, and iterate on the remaining classes that do not conflict with this selection.





## Minimum-Conflict

**Minimum-Conflict.** Select the class that conflicts with the fewest number of classes, and iterate on the classes that do not conflict with this selection.

4	4		
4	4		
4	2 i		
3	4	4	3

---

Time →



## Last-Start

**Last-Start.** Select the class that starts last, and iterate on the classes that do not conflict with this selection.

- It is not easy to find a **bad example** for this greedy algorithm.
- Because there are no bad examples! It outputs an *optimal* schedule.
- This algorithm is *symmetric* to the following greedy algorithm:  
**First-Finish.** Select the class that finishes first, and iterate on the classes that do not conflict with this selection.
- So let's prove the First-Finish greedy algorithm does indeed work...



## First-Finish Greedy Scheduling

**FirstFinish(I)**

Let Class 1 be the class with the earliest finish time.

Let  $X$  be the set of classes that clash with class 1.

Output  $\{1\} \cup \text{FirstFinish}(I \setminus X)$



# First-Finish Greedy Scheduling (Revisted)

## First-Finish Scheduling Algorithm

Sort the classes by finish times  $f_1 \leq f_2 \leq \dots \leq f_n$

Initialize the set of selected classes:  $S = \emptyset$

Initialize the set of undecided class requests:  $I = \{1, 2, \dots, n\}$

While  $I \neq \emptyset$

    Let  $i$  be the lowest index class in  $I$

    Set  $S \leftarrow S \cup \{i\}$  and  $I \leftarrow I \setminus \{i\}$

    For each class  $j \in I$

        If  $s_j < f_i$  then set  $I \leftarrow I \setminus \{j\}$

Output  $S$

### First-Finish Scheduling Algorithm

Sort the classes by finish times  $f_1 \leq f_2 \leq \dots \leq f_n$

Initialize the set of selected classes:  $S = \emptyset$

Initialize the set of open class requests:  $I = \{1, 2, \dots, n\}$

While  $I \neq \emptyset$

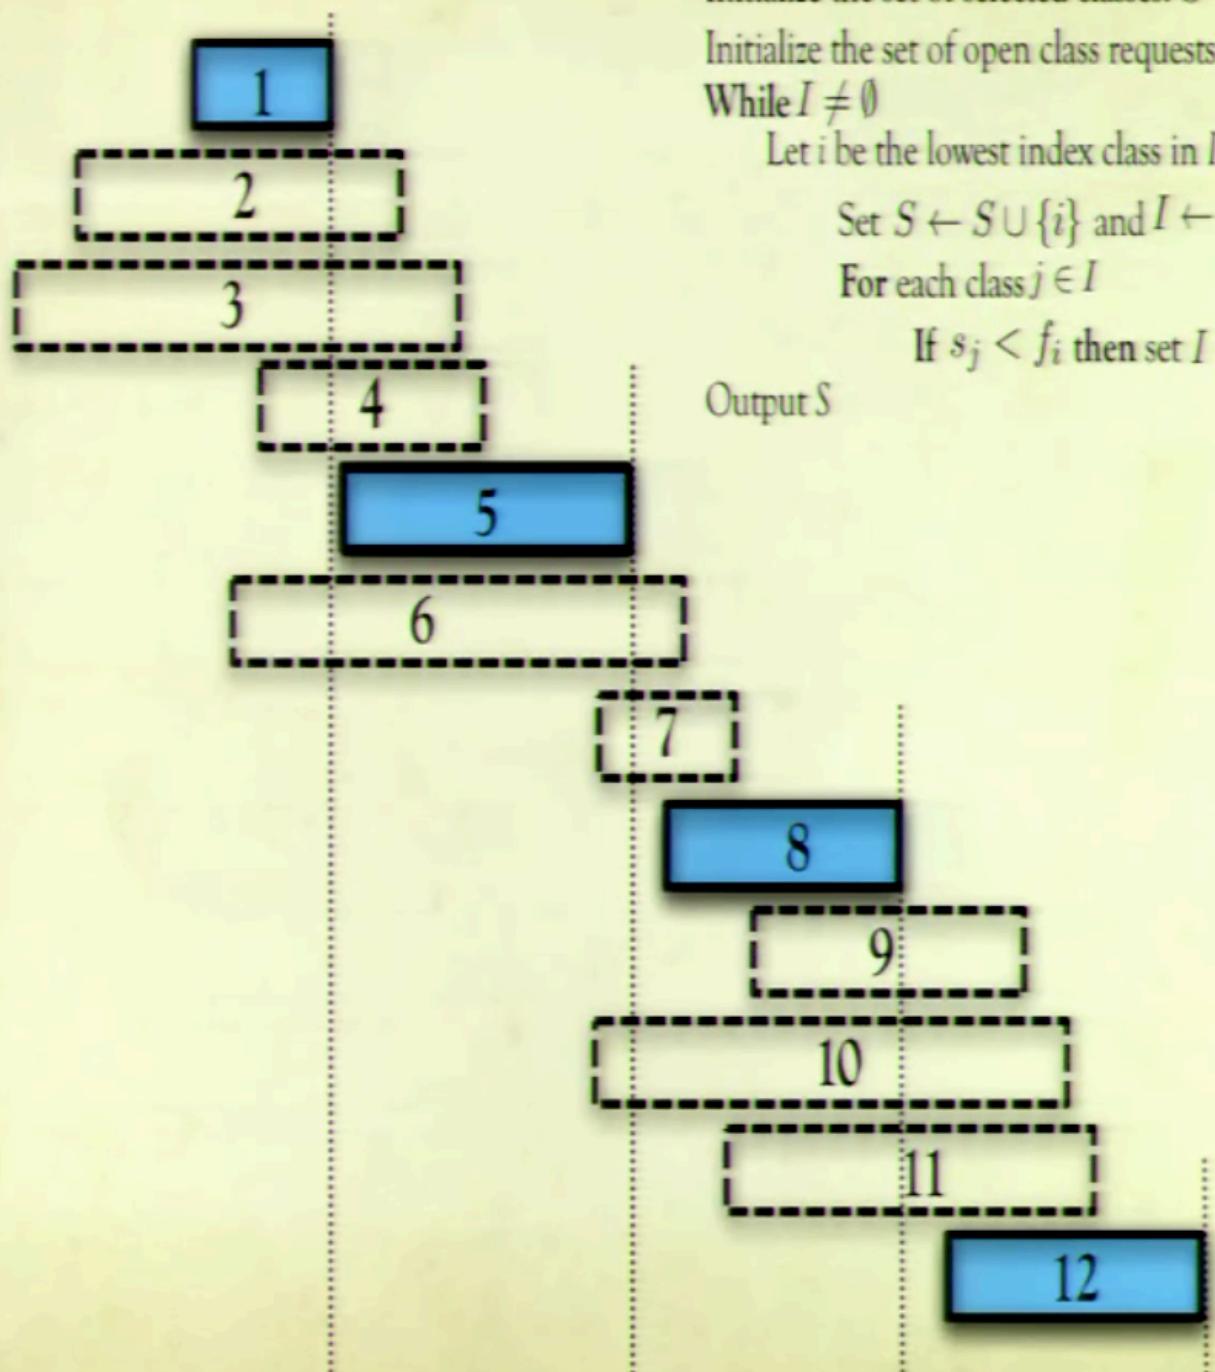
    Let  $i$  be the lowest index class in  $I$

    Set  $S \leftarrow S \cup \{i\}$  and  $I \leftarrow I \setminus \{i\}$

    For each class  $j \in I$

        If  $s_j < f_i$  then set  $I \leftarrow I \setminus \{j\}$

Output  $S$



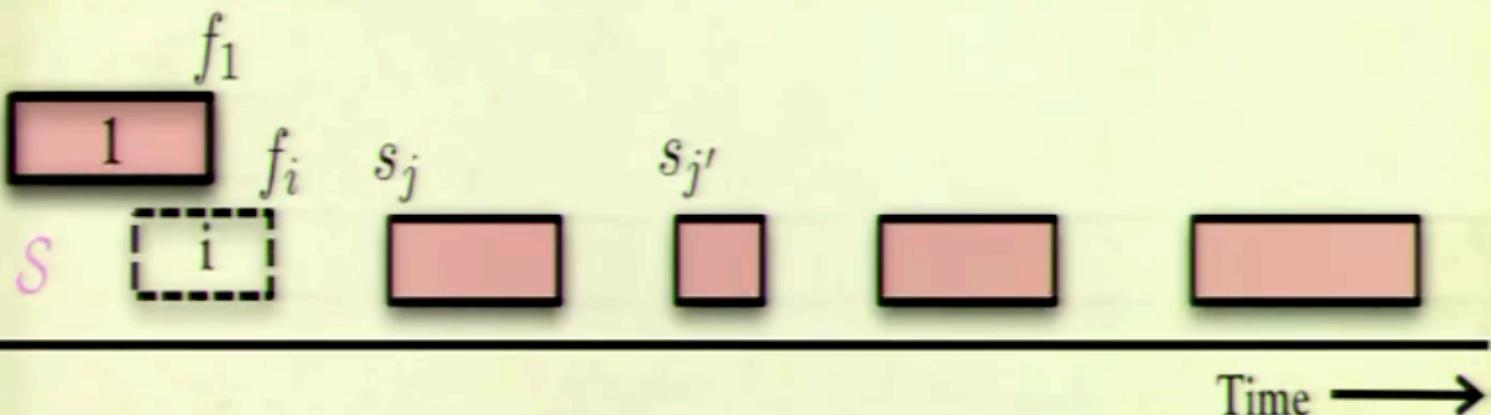
Time →

# Another Exchange Argument

**Lemma.** There is an optimal schedule that selects Class 1.

**Proof.**

- Recall the classes are indexed such that  $f_1 \leq f_2 \leq \dots \leq f_n$
- Take an *optimal schedule*  $\mathcal{S}$  and assume  $1 \notin \mathcal{S}$
- Let  $i$  be the lowest index class selected in  $\mathcal{S}$
- We claim  $\mathcal{S} \setminus \{i\} \cup \{1\}$  is a feasible allocation of maximum size.



- This follows as  $f_i \leq s_j$  for any  $j \in \mathcal{S} \setminus \{i\}$ .
- But  $f_1 \leq f_i \leq s_j$  so Class 1 does not conflict with any Class in  $\mathcal{S} \setminus \{i\}$   
 $\Rightarrow \mathcal{S} \setminus \{i\} \cup \{1\}$  is feasible with cardinality equal to  $|\mathcal{S}|$ . □

**FirstFinish( $I$ )**

Let Class  $I$  be the class with the earliest finish time.

Let  $X$  be the set of classes that clash with class  $I$ .

Output  $\{I\} \cup \text{FirstFinish}(I \setminus X)$

# The Optimality of First-Finish

**Theorem.** The first-finish algorithm outputs an *optimal* schedule.

**Proof.**

- By induction on the cardinality of the optimal solution  $|\text{opt}(I)|$

Base Case:

- Let  $|\text{opt}(I)| = 1$
- Then First-Finish outputs  $\{I\}$  and this is an optimal solution.

Induction Hypothesis:

- If  $|\text{opt}(I)| = k$  then First-Finish outputs an optimal solution.

**FirstFinish( $I$ )**

Let Class 1 be the class with the earliest finish time.

Let  $X$  be the set of classes that clash with class 1.

Output  $\{1\} \cup \text{FirstFinish}(I \setminus X)$

**Proof [cont.]**

Induction Step:

- Let  $|\text{opt}(I)| = k + 1$
- First-Finish outputs  $\{1\} \cup \text{FirstFinish}(I \setminus K)$
- But, by the lemma,  $\{1\} \in \mathcal{S}^*$  for some optimal solution  $\mathcal{S}^*$ 
  - $\implies$  Thus  $\mathcal{S}^* \setminus \{1\}$  is an optimal solution for the sub-problem  $I \setminus X$
  - $\implies |\text{opt}(I \setminus X)| = k$
- By the induction hypothesis, First-Finish gives an optimal solution for the sub-problem  $I \setminus X$ 
  - $\implies |\text{FirstFinish}(I \setminus X)| = k$
  - $\implies |\{1\} \cup \text{FirstFinish}(I \setminus X)| = k + 1$
- Thus First-Finish outputs an optimal solution. □

**FirstFinish( $I$ )**

Let Class 1 be the class with the earliest finish time.

Let  $X$  be the set of classes that clash with class 1.

Output  $\{1\} \cup \text{FirstFinish}(I \setminus X)$



## The Running Time

- How long does the first-finish scheduling algorithm take?
  - There are at most  $n$  iterations.
  - It takes  $O(n)$  time to find class that finishes earliest in each iteration.
$$\implies \text{Runtime} = O(n^2)$$
- But this is a very *crude* analysis.
- With a more subtle implementation and analysis we can obtain a running time of  $O(n \cdot \log n)$ .

# Practical Aspects



- Scheduling problems are *ubiquitous* in practice.
- However, the scheduling problems seen today should be viewed as toy examples used to introduce the topic of greedy algorithms.
- For example, for more realistic applications of interval scheduling we would want to allow for:
  - *Multiple Rooms.*
  - *Weighted Intervals.*
  - *Online Booking Requests, etc.*