

The Shortest Path Problem

Lecture 11



The Shortest Path Problem

- Given a directed graph $G=(V,A)$ with a source vertex $s \in V$.
- Let each arc $a \in A$ have a non-negative length ℓ_a
- The length of a path P in the graph is:

$$\ell(P) = \sum_{a \in P} \ell_a$$

- This induces the following problem:

The Shortest Path Problem. Find the shortest length paths from s to every other vertex $v \in V$



Dijkstra's Shortest Path Algorithm

Dijkstra's Shortest Path Algorithm [Dijkstra 1957]

Set $\mathcal{S} = \emptyset$

Set $d(s) = 0$ and $d(v) = \infty$, $\forall v \in V \setminus \{s\}$

While $\mathcal{S} \neq V$

If $v = \operatorname{argmin}_{u \in V \setminus \mathcal{S}} d(u)$ then set $\mathcal{S} \leftarrow \mathcal{S} \cup \{v\}$

For each arc (v, w) :

If $d(w) > d(v) + \ell_{vw}$ and $w \in V \setminus \mathcal{S}$

then set $d(w) = d(v) + \ell_{vw}$

Dijkstra's Shortest Path Algorithm

Set $\mathcal{S} = \emptyset$

Set $d(s) = 0$ and $d(v) = \infty$, $\forall v \in V \setminus \{s\}$

While $\mathcal{S} \neq V$

If $v = \operatorname{argmin}_{u \in V \setminus \mathcal{S}} d(u)$ then set $\mathcal{S} \leftarrow \mathcal{S} \cup \{v\}$

For each arc (v, w) :

If $d(w) > d(v) + \ell_{vw}$ and $w \in V \setminus \mathcal{S}$

then set $d(w) = d(v) + \ell_{vw}$



- This is a **greedy algorithm** as at each step it myopically selects the vertex with the smallest label.
- As described, the algorithm finds just the shortest path *distances*.
- It is easy to enhance the algorithm to keep track of the shortest paths themselves as well...

Dijkstra's Shortest Path Algorithm

Dijkstra's Shortest Path Algorithm [Dijkstra 1957]

Set $\mathcal{S} = \emptyset$

Set $\mathcal{T} = \emptyset$, $\text{pred}(s) \leftarrow “-”$, and $\text{pred}(v) \leftarrow \square$

Set $d(s) = 0$ and $d(v) = \infty$, $\forall v \in V \setminus \{s\}$

While $\mathcal{S} \neq V$

If $v = \operatorname{argmin}_{u \in V \setminus \mathcal{S}} d(u)$

then set $\mathcal{S} \leftarrow \mathcal{S} \cup \{v\}$ and $\mathcal{T} \leftarrow \mathcal{T} \cup (\text{pred}(v), v)$

For each arc (v, w) :

If $d(w) > d(v) + \ell_{vw}$ and $w \in V \setminus \mathcal{S}$

then set $d(w) = d(v) + \ell_{vw}$

$\text{pred}(w) \leftarrow v$



Dijkstra's Shortest Path Algorithm

Set $S = \emptyset$

Set $d(s) = 0$ and $d(v) = \infty, \forall v \in V \setminus \{s\}$

While $S \neq V$

If $v = \operatorname{argmin}_{u \in V \setminus S} d(u)$ then set $S \leftarrow S \cup \{v\}$

For each arc (v, w) :

If $d(w) > d(v) + \ell_{vw}$ and $w \in V \setminus S$
then set $d(w) = d(v) + \ell_{vw}$

Key



Vertex (Unselected)

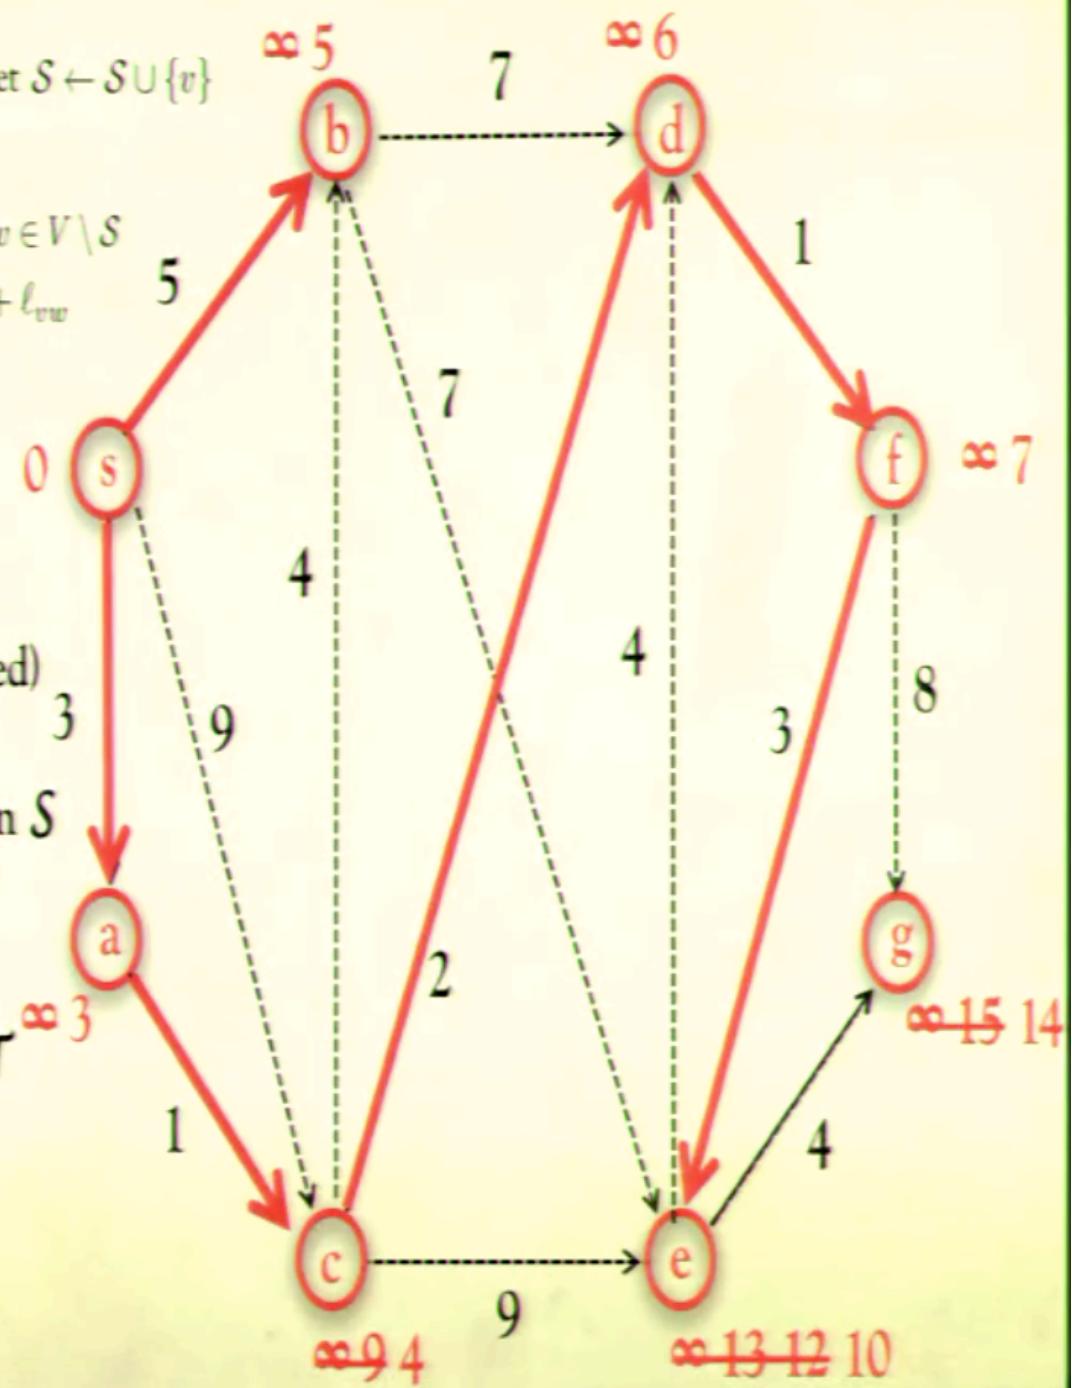


Selected Vertex in S

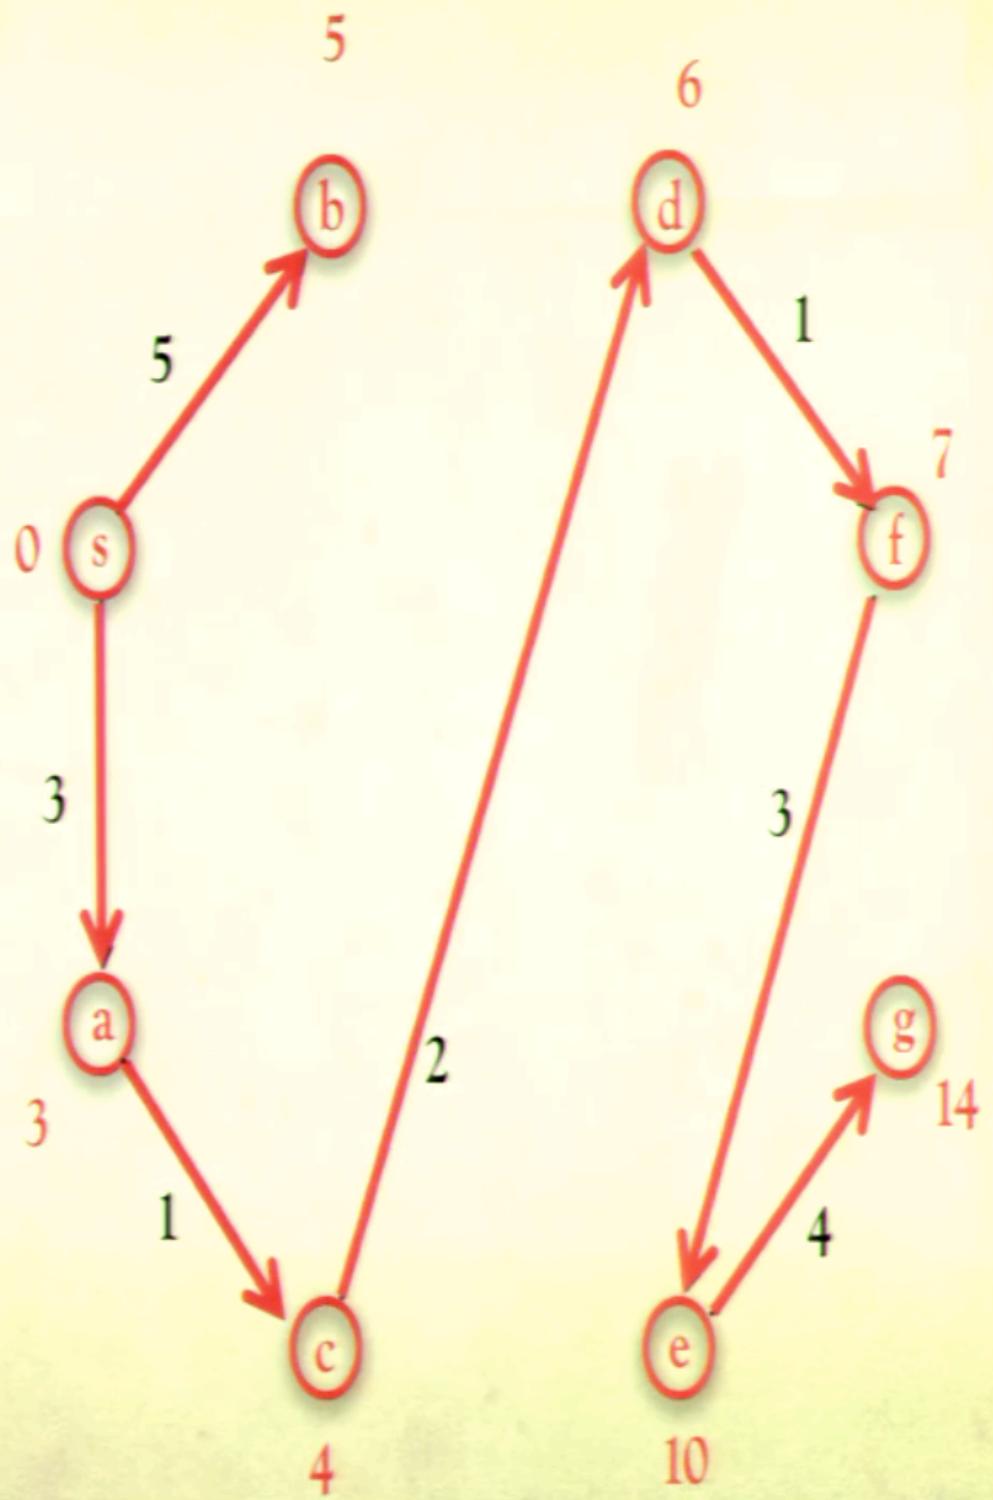


Predecessor Arc

Selected Arc in T

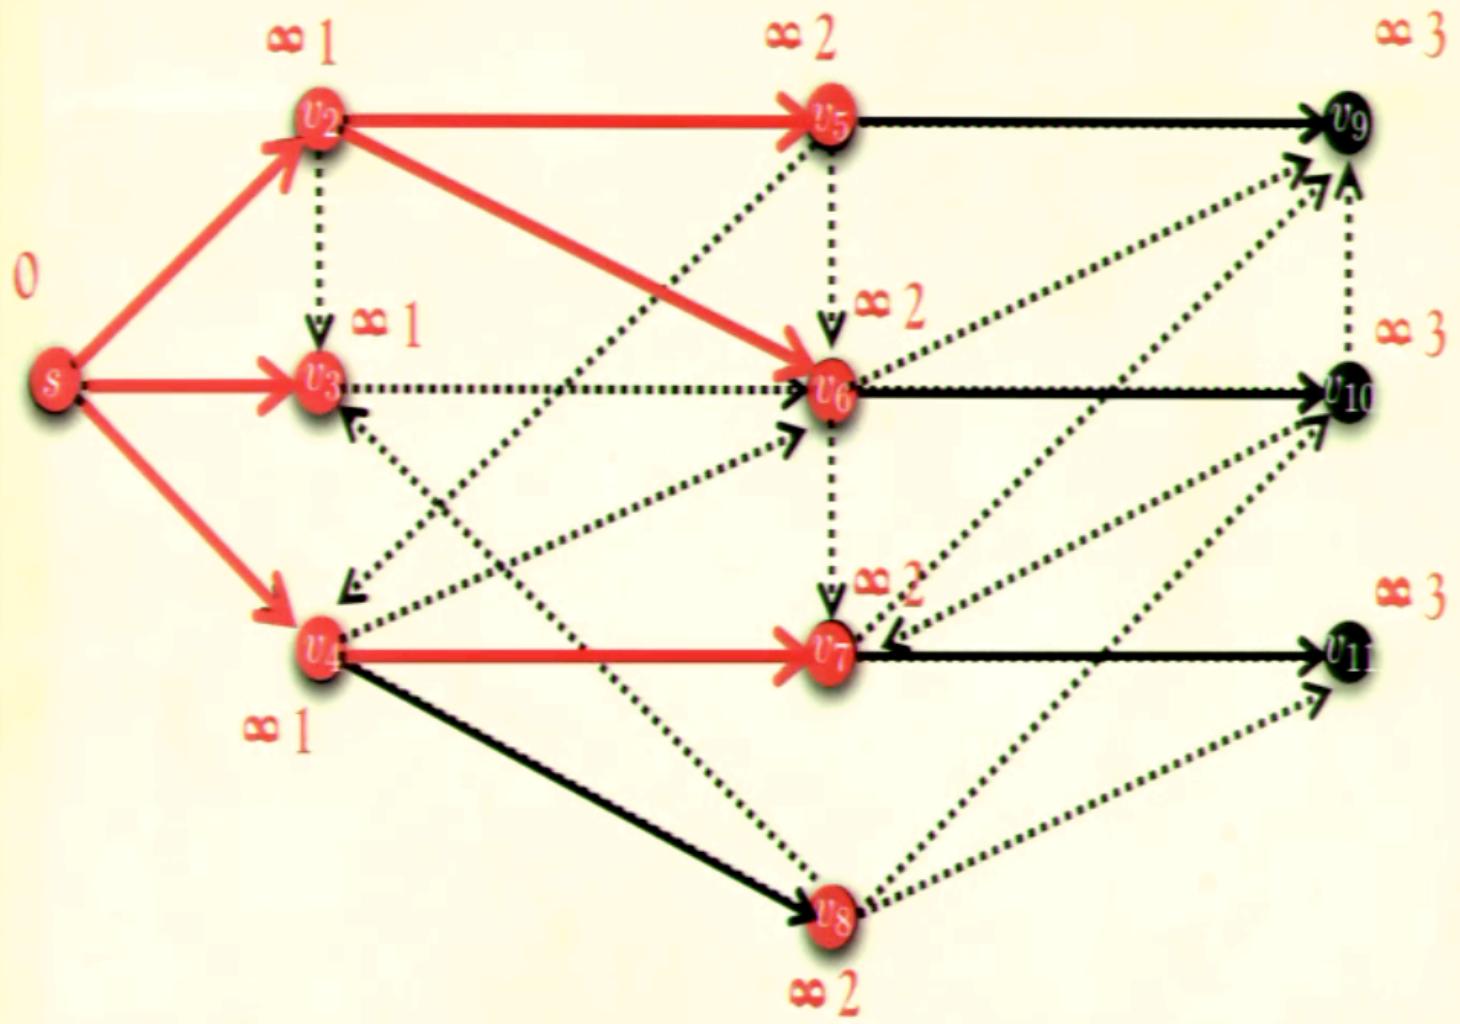


The Shortest Paths





What happens in the special case
where each arc length is exactly one?



Key



Vertex (Unselected)



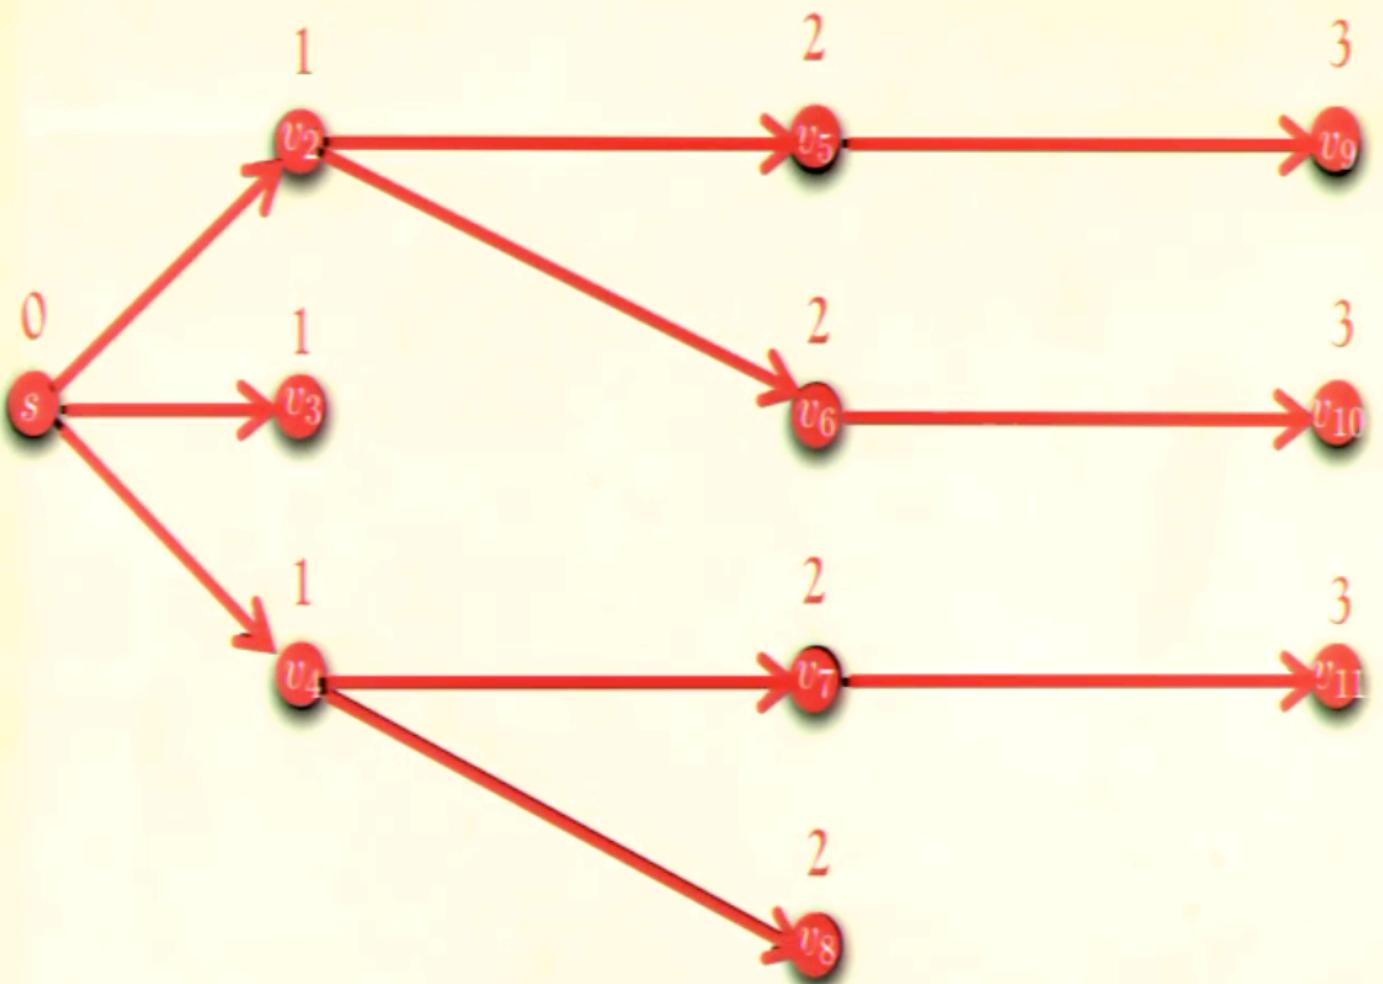
Predecessor Arc

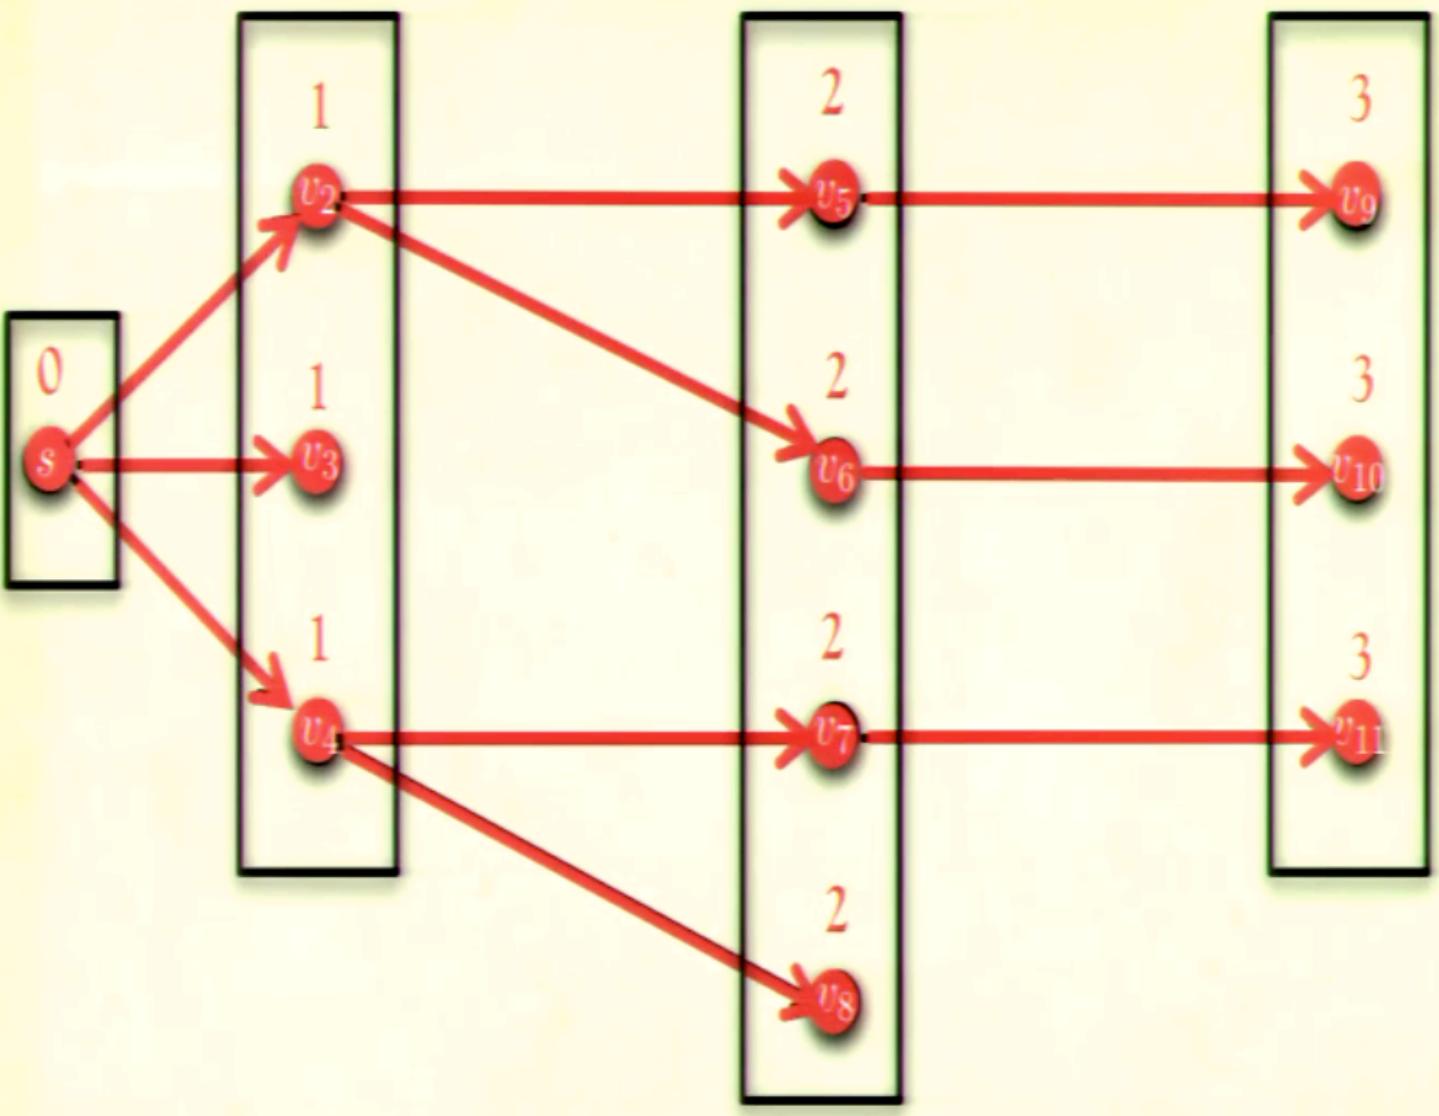


Selected Vertex in S



Selected Arc in T





- This is the breadth-first search algorithm!

The Shortest Path Graph

- To show the algorithm does work, first observe that S is a set of vertices and T is a set of arcs.
- The following *notation* will then be helpful:
 - Let S^k be the set of vertices in S at the end of the k th iteration.
 - Let T^k be the set of arcs in T at the end of the k th iteration.
- Furthermore, observe the arcs in T^k are all between vertices in S^k .
 $\implies \mathcal{G}^k = (S^k, T^k)$ is a directed graph.
- Finally, observe that $\mathcal{G} = \mathcal{G}^n$ is the final output of Dijkstra's algorithm.
- So let's investigate what \mathcal{G} looks like...



The Shortest Path Tree

- It turns out that $\mathcal{G} = \mathcal{G}^n$ is a directed tree (*arborescence*) rooted at s .
- To show this we prove:

Theorem. The graph \mathcal{G}^k is a directed tree (*arborescence*) rooted at s .

Proof.

Base Case: $k=1$

- Observe S^1 contains only the vertex s .
- But T^1 contains no arcs as $pred(s) = “-”$
- Thus $\mathcal{G}^1 = (S^1, T^1)$ is trivially an *arborescence* rooted at s .

Induction Hypothesis: $\mathcal{G}^{k-1} = (S^{k-1}, T^{k-1})$ is an arborescence rooted at s .

Proof [cont].

Induction Step: Consider $\mathcal{G}^k = (\mathcal{S}^k, \mathcal{T}^k)$

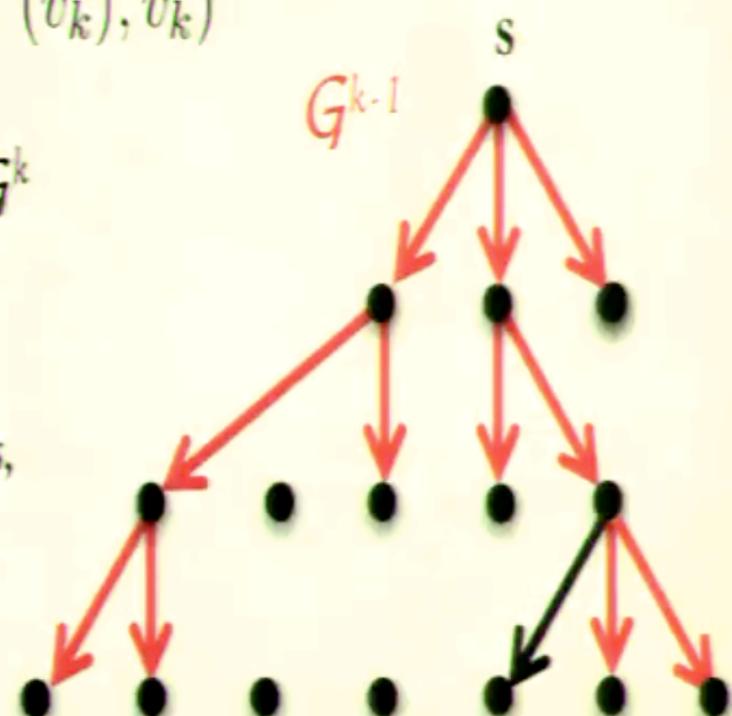
- Let v_k be the vertex added to \mathcal{S} in the k th iteration.

$$\Rightarrow \mathcal{S}^k = \mathcal{S}^{k-1} \cup \{v_k\}$$

$$\Rightarrow \mathcal{T}^k = \mathcal{T}^{k-1} \cup (\text{pred}^{k-1}(v_k), v_k)$$

- Thus v_k has in-degree exactly 1 in G^k and out-degree exactly 0.

- Moreover, by the *induction hypothesis*, G^{k-1} is an arborescence rooted at s .



- Furthermore, by definition: $\text{pred}^{k-1}(v_k) \in \mathcal{S}^{k-1}$

$\Rightarrow \mathcal{G}^k = (\mathcal{S}^k, \mathcal{T}^k)$ is an arborescence rooted at s with v_k as a leaf. □



Shortest Path Distances

- So Dijkstra's algorithm outputs an arborescence.
- Moreover the paths induced by the arborescence are shortest paths...

Theorem. The graph $\mathcal{G}^k = (\mathcal{S}^k, \mathcal{T}^k)$ gives the true shortest path distances from s to every vertex in \mathcal{S}^k .

Proof.

Base Case: $k=1$

- Trivially true as: $d^1(s) = 0 = d^*(s)$

Induction Hypothesis:

- Assume that $\mathcal{G}^{k-1} = (\mathcal{S}^{k-1}, \mathcal{T}^{k-1})$ gives the shortest distance from s to *every* vertex in \mathcal{S}^{k-1} :

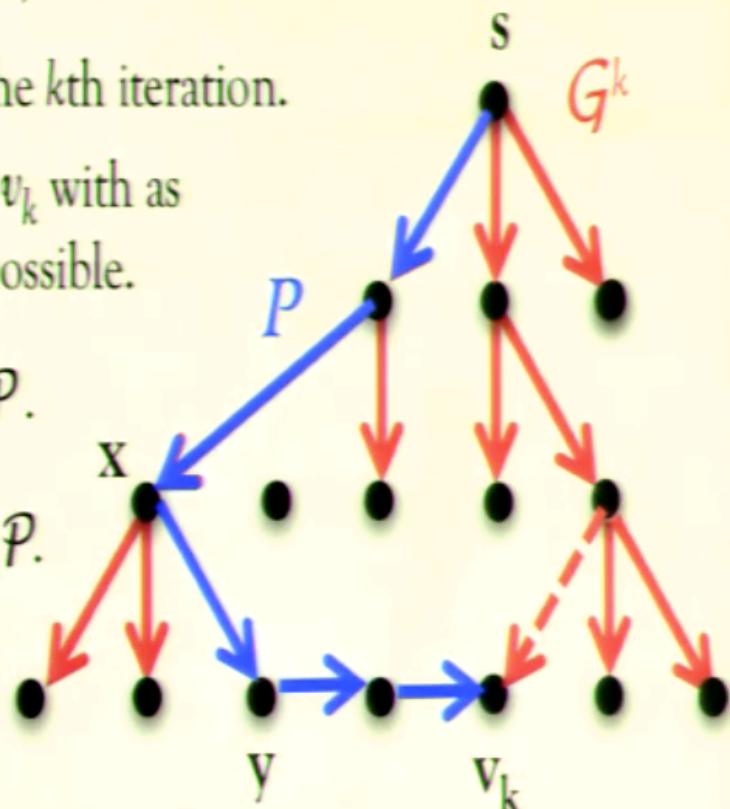
$$d^{k-1}(v) = d^*(v) \quad \forall v \in \mathcal{S}^{k-1}$$

Proof [cont].

Induction Step: Consider $\mathcal{G}^k = (S^k, T^k)$

- Let v_k be the vertex added to S in the k th iteration.
- Take the shortest path P from s to v_k with as many arcs in common with G^k as possible.
- Let x be the last vertex of G^{k-1} in P .
- Let $y \notin S^k$ be the vertex after x in P .

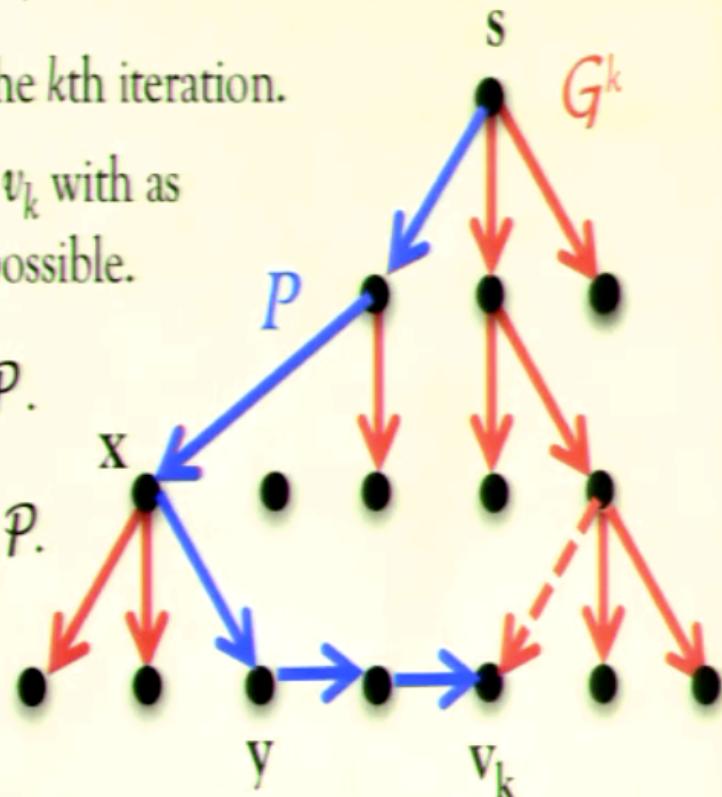
Such a y exists or P is in G^k and we are done.



Proof [cont].

Induction Step: Consider $\mathcal{G}^k = (\mathcal{S}^k, \mathcal{T}^k)$

- Let v_k be the vertex added to \mathcal{S} in the k th iteration.
- Take the shortest path P from s to v_k with as many arcs in common with G^k as possible.
- Let x be the last vertex of G^{k-1} in P .
- Let $y \notin \mathcal{S}^k$ be the vertex after x in P .



- Since y is on the shortest path from s to v_k and the arc lengths are non-negative we have:

$$d^*(y) \leq d^*(v_k)$$

$$< d^k(v_k) -$$

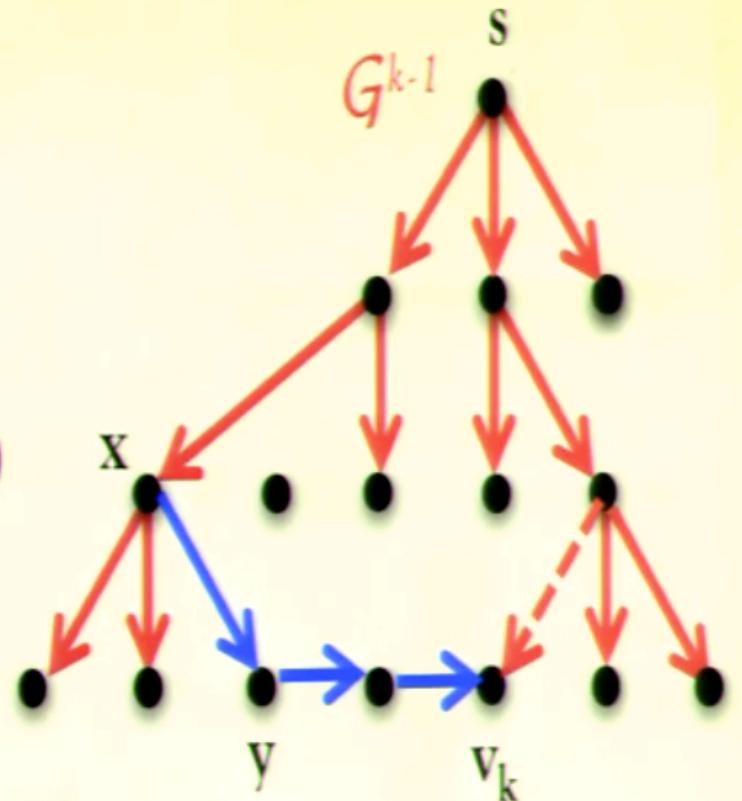
We have a strict inequality or we are done.

Proof [cont].

- So we have: $d^*(y) < d^k(v_k)$

- But since $x \in S^{k-1}$ we have:

$$\begin{aligned}
 d^k(y) &\leq d^{k-1}(x) + \ell(x, y) \\
 &= d^*(x) + \ell(x, y) \\
 &= d^*(y)
 \end{aligned}$$



- Therefore:

$$d^k(y) \leq d^*(y) < d^k(v_k)$$

- This contradicts the selection of v_k as: $d^k(y) < d^k(v_k)$





The Running Time

- How long does Dijkstra's algorithm take?
 - There are n iterations.
 - There are at most n distance updates in each iteration.
$$\Rightarrow \text{Runtime} = O(n^2)$$
- In fact, later in the course, we will see how to implement Dijkstra's algorithm in time $O(m \cdot \log n)$ using a data structure called a heap.

Applications of Dijkstra's Algorithm

- Despite being 60 years old, Dijkstra's algorithm and its variants are still widely used in practice. For example in:
 - Open Shortest Path First (OSPF), the most common *Interior Gateway Protocol* (IGP) for routing internet packets in an autonomous system.
 - Route Mapping

